Stephen Chung

Nottingham, United Kingdom | +44 7884 768773 | chung25stephen@gmail.com | linkedin.com/in/stephen-chung

EDUCATION

The University of Nottingham   **Nottingham, United Kingdom**
MSc Computer Science (Artificial Intelligence) – Ongoing   *September 2024 – September 2025*
- Dissertation: *Agent-Based LLMs for Analysing Complex Features in Videos and Images*

University of California, Davis   **Davis, CA, USA**
BA in Economics & Minor in Computer Science   *September 2016 – July 2020*

SKILLS

Frontend/Mobile: Javascript; Typescript; Dart; Next.JS; React; AngularJS; Flutter
Backend: Python; Golang; C#; Flask; FastAPI; Go Gin; .NET
Databases: PostgreSQL; MySQL; MongoDB; DynamoDB; Neo4j
LLM: LangChain; Atomic Agents; LlamaIndex; Hugging Face; Milvus DB; Qdrant
Certificate: AWS Certified Solutions Architect – Associate (issued by AWS in April 2021)

EXPERIENCE

Kiitzu (Software House)   **Taipei, Taiwan**
Full Stack Engineer   *March – September 2024*
- Architected and delivered a bilingual learning mobile application using Flutter, increasing user retention by 25% through seamless authentication with Firebase and targeted user insights via Google Analytics (live on Google PlayStore: One-Forty).
- Developed responsive web front-ends with Next.JS, ShadCn, and TailwindCSS, implementing JavaScript's Intersection Observer API to create scroll animations that improved user engagement metrics.
- Transformed company's AWS infrastructure from EC2 instances to ECS clusters, reducing monthly cloud costs by 30% while improving system scalability and deployment reliability.

Advanced Micro Devices (AMD)   **Taipei Taiwan**
Full Stack Engineer   *June 2023 - Jan 2024*
- Spearheaded the front-end development of an automated GPU testing web interface for AMD's Navi Series GPUs, delivering a Next.JS solution with Material UI and Zustland to accelerate product testing cycles by 50%.
- Architected and deployed the cloud infrastructure essential for application deployment, enabling seamless scalability and reliability that expanded the user base to 2 other regional offices.
- Developed a Golang daemon service with GORM and Gin Web Framework to provide real-time monitoring for Windows test servers, enabling remote control for over 40+ Windows machines.

Trend Micro   **Taipei, Taiwan**
Software Engineer   *September 2021 - June 2023*
- Delivered and maintained critical components of Trend Micro's business process management system handling 10,000+ monthly tickets, using a full-stack approach with C# .NET, Python and AngularJS.
- Designed and built a cross-regional OCR invoice scanning system using AWS Step Functions and Azure Form Recognizer, reducing manual data entry by 80%.
- Adopted Neo4j graph database to replace legacy permission routing systems, reducing query time by 25% while streamlining complex organizational hierarchy management and approval workflows.

SELLERLINX (Ecommerce SaaS Startup)   **Taipei, Taiwan**
Backend Engineer   *September 2020 – September 2021*
- Built and maintained multiple core systems that directly enabled the successful platform launch, scaling to 200+ active business users within the first quarter.
- Developed AWS Step Functions-powered bulk listing feature for Shopee merchants, enabling 200+ product management simultaneously, saving hours of manual inventory management for clients.
- Created Lambda/SQS automated billing system that processed platform-wide invoices monthly, replacing 100% of manual platform fee accounting.

PROJECTS

Skincare GPT (https://github.com/spchung/skincare_gpt)   **Personal Project**
- Developed RAG application analyzing Sephora skincare database with OpenAI embeddings pipeline, enabling personalized product recommendations based on user needs.
- Implemented custom intent classification system with Qdrant vector database schema, optimizing search accuracy and query speed for different response types.

Multi-Agent Debate (https://github.com/spchung/agent_debate)   **University Project**
- Designed LLM-powered platform enabling structured academic debates between AI agents with varying access to shared research knowledge base.
- Engineered Planning Agent architecture to extract and strategically select key arguments, alongside a RAG-powered Retrieval Agent for contextual responses.
- Developed Knowledge Graph Agent generating interconnected argument nodes with repetition prevention, with LogFire monitoring for performance analysis.